

Link Prediction on a Multi-Relational Knowledge Graph: FB15k-237, Matched KGE, and Planned Hard-Negative & Slice Analysis

Thalia Delhaise, Lenny Briclet, André Fonseca, Rafael Gufler (DL + KG joint project)

RC2526 — Deliverable 2 (Report part 1)

1. Introduction and motivation

A **knowledge graph** (KG) is a set of binary facts (h, r, t) over a finite set of **entities** $\mathcal{E}$ and **relation types** $\mathcal{R}$. **Link prediction** (knowledge base completion) asks which missing triples are true: for a query with one side hidden, e.g. $(h, r, ?)$, rank all candidate tail entities. This is a standard **learning on graphs** task with applications to semantic query answering and maintenance of large knowledge bases.

In this project we (i) **materialize** a benchmark KG from the **FB15k-237** public triple files, (ii) run **TransE** and **ROTATE** under matched conditions, and (iii) develop **custom hard-negative** training and **relation-wise** evaluation—the main methodological extensions. A **relational GNN** is a **optional** add-on for the final phase if time allows; it is not required for the core claim.

Our **contribution** is: reproducible KGE baselines; **own PyTorch** code for harder negatives and ablations vs. defaults; and **analysis** (slices, error cases) on the **same filtered** metrics (Hits@ k , MRR).

2. Problem statement

Formal graph. $\mathcal{G} = (\mathcal{E}, \mathcal{R}, \mathcal{T})$ with $\mathcal{T} \subseteq \mathcal{E} \times \mathcal{R} \times \mathcal{E}$. The graph is **directed** and **multi-relational**: the same node pair can appear in several relations. Training sees only $\mathcal{T}_{\text{train}}$. Validation and test are disjoint held-out **positive** triples.

Data. FB15k-237 (Freebase-derived) [4]: 14,541 entities, 237 relations, 272,115 / 17,535 / 20,466 training, validation, and test triples. The benchmark is widely used; inverse relations are removed to reduce trivial inferences. We load the official split via **PyKEEN** [1]; optional RDF/edge export for a KG visualization slide is a minor technical step.

Task. Link prediction: for each test triple, rank the correct head or tail under the **filtered** setting [2] (candidates that complete a *known* triple in train, val, or test are not errors).

Success criteria. (1) Strong **aggregate** filtered MRR and Hits@1,3,10. (2) Ablations on **negative** strategies. (3) Deeper **relation-wise** and **degree** analysis and short error study. (4) **Optional:** a GNN baseline. We do not claim SOTA; we care **where** each setting helps or fails.

3. Methodology

Phase A — Classical KGE (Milestone 1–2). We implement **TRANSE** and **ROTATE** [4] with **PyKEEN** for a fair **first** line: same splits, d , optimiser, training budget, and **filtered** evaluation. The scoring function is learned; training uses margin/ranking loss and negative sampling. This serves as a **baseline** for the Deep Learning report.

Phase B — Custom negative mining (in-house PyTorch). In parallel with KGE, we will implement a **learnt or scheduled hard-negative** mechanism: e.g. a small g_ϕ over a subset of corruptions, or periodic **mining** in embedding space, with ablations **against** uniform and Bernoulli **defaults**. This addresses the “not only default PyKEEN” feedback.

Phase C (optional) — Relational GNN. If time allows, a shallow R-GCN (or COMPGCN) in PyG/DGL [3] on the **same** task and splits; not required for the main thesis (hard negatives + KGE + analysis).

Evaluation. In addition to global MRR and Hits@ k , we will report **slices** (relation frequency, entity degree) and a short **error** / case study. Optional bootstrap confidence bands on the test set.




Figure 1: Core pipeline: same filtered evaluation; custom negatives and analysis on the KGE line. GNN (not shown) is optional for later.

4. Planned experiments and preliminary results

Planned. (1) TransE/RotatE; **match** d and budget. (2) Custom negative ablations. (3) Slice + error analysis. (4) Optional GNN. (5) Time/hardware.

Preliminary (Milestone 1). We use the same training script and protocol as the Deep Learning M1 report (`code/train_baseline_kge.py`): $d = 128$, batch 1024, Adam lr 10^{-3} , up to 15 epochs with early stopping (patience 3) on validation MRR. Test metrics in Table 1 are aligned with that report; PyKEEN reports **realistic** filtered ranking, head and tail prediction averaged (**both** in PyKEEN’s evaluation).

Table 1: Preliminary aggregate test results, FB15k-237 (same numbers as DL Milestone 1; realistic filtered MRR / H@10, **both**).

Setting	MRR (filtered)	Hits@10 (filtered)
TransE	0.129	0.263
RotatE	0.234	0.378
GNN (optional)	—	—

5. Acknowledgment of course deliverables (D2 scope)

D2 (this document) is the **first** report slice: **introduction, problem, methodology, and planned/early results**, to be extended in D3 with full experiments, student contributions, and the Jupyter lab notebook, per RC2526 project text.

References

- [1] Mehdi Ali, Max Berrendorf, Charles Tapley Hoyt, Laurent Vermue, et al. PyKEEN 1.0: A python library for training and evaluating knowledge graph embeddings. *JMLR*, 22(82):1–6, 2021.
- [2] Antoine Bordes, Nicolas Usunier, Alberto Garcia-Durán, Jason Weston, and Oksana Yakhnenko. Translating embeddings for modeling multi-relational data. In *NeurIPS*, 2013.
- [3] Michael Schlichtkrull, Thomas N Kipf, et al. Modeling relational data with graph convolutional networks. In *ESWC*, 2018.
- [4] Zhiqing Sun, Zhi-Hong Deng, Jian-Yun Nie, and Jian Tang. RotatE: Knowledge graph embedding by relational rotation in complex space. In *ICLR*, 2019.